

Week 1: Introduction & Computational Models

Instructor: Warut Suksompong

National University of Singapore

CS3230


Semester 2, 2021–22

Introduction

- Module objectives:
 - To study algorithms in a formal way (through the lens of mathematics)
 - To learn tools to analyze the performance of algorithms
 - To learn techniques to design an efficient algorithm
- After the module, students should be able to
 - Perform analysis of the asymptotic performance of algorithms.
 - Design efficient algorithms to solve problems.
 - Able to comment on correctness of designed algorithms.
 - Comment on (inherent) hardness of a problem.

Introduction

- Prerequisites:
 - CS2010 or CS2020 or CS2040/C/S Data Structures and Algorithms
 - CS1231/S or MA1100 Discrete Structures
- Textbook:
 - **CLRS**: Introduction to Algorithms, 3rd edition, by Cormen, Leiserson, Rivest & Stein, 2009



Logistics

- **Lecture:** Thursday 14:00–16:00 via Zoom, recorded
 - I will stick around to answer questions after lecture.
 - Lecture notes and slides will be posted on LumiNUS.
 - My email address: warut@comp.nus.edu.sg
 - Due to the large number of students, please don't email me questions about course material. Post them on the LumiNUS forum instead.
- **Tutorial:** 1 hour each week, starting from week 3
 - 18 slots (10 F2F, 8 via Zoom)
 - Schedule and list of tutors in “Tutorials” folder on LumiNUS (soon).
 - Each tutor will also hold one hour of office hours per week.
- **LumiNUS forum:** Ask questions here!
 - I will monitor it together with TAs.
 - The posts will show your “nickname”. You can set your nickname in your LumiNUS user profile.

Tentative Schedule

Week	Date	Topic
1	13 Jan	Intro & computational models
2	20 Jan	Asymptotic analysis
3	27 Jan	Iteration, recursion & divide-and-conquer
4	3 Feb	Average-case analysis & randomized algorithms
5	10 Feb	Hashing
6	17 Feb	Pattern matching & Streaming
7	3 Mar	Midterm (during lecture slot)
8	10 Mar	Amortized analysis
9	17 Mar	Dynamic programming
10	24 Mar	Greedy algorithms
11	31 Mar	Reductions & computational complexity
12	7 Apr	Reductions & computational complexity (cont.)
13	14 Apr	No class (NUS Well-Being Day)

Assessment

- Assignments (36%):
 - 12 assignments, one for each lecture (last assignment will be a revision)
 - Each assignment consists of 3 questions.
 - There is one question graded for correctness in every odd-numbered assignment (1, 3, 5, 7, 9, 11). All other questions are graded for effort.
 - Each correctness-based question worth 3.5% (graded out of 7 points)
 - Each effort-based question worth 0.5% (graded out of 1 point)
 - Total = $6 \cdot 3.5\% + 30 \cdot 0.5\% = 36\%$
- Assignment released on lecture day (Thursday), due 11:59pm Sunday of the following week (except Assignment 12)
- Assignment schedule + grader list posted on LumiNUS (in the “Assignments” folder)
- **No late assignment will be accepted.**
- For grading enquiries, please check directly with the relevant grader.

Assessment

- Continuous assessment (40%):
 - Assignments (36%)
 - Tutorial attendance (4%)
 - Bonus points (up to 6%)
 - Free for everyone! (2%)
 - Two programming assignments (total of 4%)
 - The total you earn from continuous assessment cannot exceed 40%.
- Exams (60%):
 - Midterm (30%): **3 March** (during lecture slot), **14:00–16:00**
 - Final exam (30%): **26 April** (Tuesday), **17:00–19:00**
 - Mode of both exams will be online.

Problems v Algorithms

- Problems provide the **what**. Algorithms provide the **how**.
- Example of a computational problem: **Multiplication**
 - **Input:** Two numbers x and y
 - **Output:** The product $x \cdot y$
- An **algorithm** is a well-defined procedure for finding a correct solution to the input.
- There can be many algorithms for a particular problem.
- The grade-school algorithm for multiplication is just one algorithm for the problem (in fact, not the best one when the input is large!)

Algorithms

- In this module, we will focus on two key aspects of algorithms: **correctness** and **efficiency** (i.e., **running time**).
- For correctness, we typically want algorithms to be correct on **every** valid input.
- This is known as **worst-case correctness**.
- Sometimes, relaxed notions of correctness are considered:
 - Correct on a random input
 - Correct on every input with high probability
 - Approximately correct

Algorithms

- The **running time** measures the number of steps executed by an algorithm as a function of the **input size**.
 - What each **step** is depends on the computational model used.
- We must specify what the notion of input size is.
 - If input is an **array**, its size is typically the length of the array.
 - If input is a **number**, its size is typically the length of its binary representation.
 - If input is a **graph**, its size is typically the number of vertices and edges in the graph.
- The **(worst-case) running time** of an algorithm is the **maximum** number of steps executed when run on an input of size n .
- Besides correctness and running time, other considerations include simplicity, space usage, energy consumption, parallelism, and fairness/ethics.

Comparison Model

- The input is an array of n numbers.

3	7	5	8	4
---	---	---	---	---

- The algorithm can **compare** any two elements in one time unit:
 $Is\ x > y, x < y, \text{ or } x = y?$
- No other operations on the elements (e.g., addition, subtraction) are allowed.
- Running time = total number of comparisons made
- The array can be manipulated (e.g., permuted or broken into subarrays) at no cost.

Comparison Model: Maximum

- **Problem:** Given an array A of n distinct elements (denoted by A_1, \dots, A_n), find the largest element in A .
- Here is one algorithm:


RUNTHRU(A)

```
1   $cur = 1$ 
2   $n = A.length$ 
3  for  $i = 2$  to  $n$ 
4      if  $A_i > A_{cur}$ 
5           $cur = i$ 
6  return  $A_{cur}$ 
```


- RUNTHRU makes exactly $n - 1$ comparisons for any input.

Maximum

- **Claim:** Every algorithm solving the Maximum problem must make $\geq n - 1$ comparisons!
- Let's look at some small cases first.
 - $n = 2$: Need to do 1 comparison.




- $n = 3$: If the algorithm makes just 1 comparison (say, between the first and second elements), then the third element is left untouched.



The third element could be either **very large** (in which case it is the maximum), or **very small** (in which case it is **not** the maximum).
So 2 comparisons are needed.



Maximum

- To prove the claim for general n , fix an algorithm \mathcal{M} that solves the Maximum problem on all inputs using $< n - 1$ comparisons.
- Take an input array A on which \mathcal{M} makes $< n - 1$ comparisons.
- Construct a graph G on n nodes v_1, \dots, v_n , where there is an edge between nodes v_i and v_j iff \mathcal{M} compares A_i and A_j .



Maximum

- Since G has $< n - 1$ edges, it is **disconnected**.



- Let A_i be the maximum element of A .
- Consider a different input A' , where all numbers in a different connected component than A_i are increased by a huge amount.
- \mathcal{M} cannot distinguish between A and A' , a contradiction.

Maximum


- This proof is an example of an **adversary argument**.
- The input is decided on-the-fly by an adversary who keeps its options open about what the actual input is.
- The adversary makes sure that if the algorithm makes too few comparisons, then:
 - There are **two different inputs** which are **consistent** with the results of these comparisons.
 - And yet, the solutions for the two inputs are **different**.

Second Largest

- What about finding the **second largest element**?
- One way is to find the maximum first using our previous algorithm, then use this algorithm **again** to find the maximum among the remaining $n - 1$ numbers.
- This requires $(n - 1) + ((n - 1) - 1) = 2n - 3$ comparisons.
- Is this the best we can do?
- **No!**
- Charles Lutwidge Dodgson (better known as Lewis Carroll, author of *Alice in Wonderland*), came up with an algorithm requiring fewer comparisons.


Second Largest

- Instead of finding the maximum using the previous algorithm, we do so by comparing elements using a **knockout tournament** structure.



- The winner of this “tournament” is the maximum element.
- Since every non-winner has “lost” exactly once, this algorithm also solves the Maximum problem using $n - 1$ comparisons.

Second Largest



- Observe that the second-largest element must have lost to the winner.
- We can therefore find the maximum among the $\lceil \lg n \rceil$ elements that lost to the winner using $\lceil \lg n \rceil - 1$ comparisons (\lg denotes \log_2).
- The total number of comparisons is therefore $(n - 1) + (\lceil \lg n \rceil - 1) = n + \lceil \lg n \rceil - 2$.
- This is known to be **optimal!**

Sorting

- For the Sorting problem, we want to order **all** elements in the array A of distinct numbers.
- For simplicity, assume that n is a power of two.
- **Claim:** There is a sorting algorithm that requires $\leq n \lg n - n + 1$ **comparisons**.
- For example, the **Merge Sort** algorithm (covered in Lecture 3).
- The algorithm divides A into two equal halves and merges them into one sorted array.
- If each half contains $n/2$ elements, the merging step takes $n - 1$ comparisons.

Sorting

- **Claim:** Every sorting algorithm must make $\geq \lg(n!)$ comparisons.
- Initially, there are $n!$ permutations of the set $\{1, \dots, n\}$ that the adversary could choose as the array A . Call this set \mathcal{U} .
- Each permutation in \mathcal{U} needs to be ordered differently to get sorted.
- When a query comes in (“Is $A_i > A_j$?”), the adversary checks whether $\mathcal{U}_{\text{yes}} = \{A \in \mathcal{U} : A_i > A_j\}$ is of size at least $|\mathcal{U}|/2$.
 - If so, it replies **Yes** to the algorithm and sets \mathcal{U} to be \mathcal{U}_{yes} .
 - Else, it replies **No** and sets \mathcal{U} to be $\mathcal{U} \setminus \mathcal{U}_{\text{yes}}$.
- If the algorithm makes $< \lg(n!)$ comparisons, \mathcal{U} will still contain at least two permutations, since its size decreases by at most half with each comparison.
- The algorithm will order these two permutations in the same way, and will be wrong on at least one of them.

- How big is $\lg(n!)$ then?
- It is known that $n! \geq (n/e)^n$, so

$$\lg(n!) \geq n \lg \left(\frac{n}{e} \right) = n \lg n - n \lg e \approx n \lg n - 1.44n.$$

- This means that roughly $n \lg n$ comparisons are both **required** and **sufficient** for sorting n numbers.

Query Model: Strings


- Suppose the input is a string of n bits (each bit is 0 or 1).
- With each query, the algorithm can find out **one bit** of the string.
- Consider the **problem** of **deciding whether an input string is the all-0 string or not**.
- By checking whether each bit is 0, **n queries suffice**.
- **Claim:** n queries are also necessary.
 - Suppose \mathcal{M} is an algorithm making $< n$ queries.
 - Consider an adversary that replies to each of \mathcal{M} 's queries with 0.
 - After \mathcal{M} halts, there is still at least one unqueried bit, say the i -th bit.
 - The input may be all-0, or it may be 0 in all bits except the i -th bit.

Query Model: Graphs


- The input is the (symmetric) adjacency matrix of an n -node undirected graph.
- With each query, the algorithm can find out one entry of the matrix (i.e., whether an edge is present between two chosen nodes or not).
- Consider the problem of deciding whether a graph is connected or not.
- **Claim:** $\binom{n}{2}$ queries are necessary.
 - Suppose \mathcal{M} is an algorithm making $< \binom{n}{2}$ queries.
 - When \mathcal{M} makes a query, the adversary tries not adding this edge, but adding all remaining unqueried edges.
 - If the resulting graph is connected, the adversary replies 0 (i.e., edge does not exist).
 - Else, the adversary replies 1 (i.e., edge exists).

Query Model: Graphs


- **Example:** Consider a graph with 4 nodes v_1, v_2, v_3, v_4 .
- First query: $(v_1, v_2) \rightarrow 0$ (edge does not exist)



- Second query: $(v_1, v_3) \rightarrow 0$ (edge does not exist)



- Third query: $(v_1, v_4) \rightarrow 1$ (edge exists)




Query Model: Graphs

- At every stage, setting all unqueried entries to 1 will make the graph connected.
- At the end, since \mathcal{M} made $< \binom{n}{2}$ queries, at least one entry of the adjacency matrix is unqueried.
- The adversary considers the graph G_0 obtained by setting all unqueried entries to 0, and the graph G_1 obtained by setting all unqueried entries to 1.
- By the first bullet point above, G_1 is connected.
- **Claim:** G_0 is disconnected.
- This claim suffices to finish the proof.

Query Model: Graphs

- **Claim:** G_0 is **disconnected**.
- Let (i, j) be an unqueried pair of nodes.
- Suppose for contradiction that there is a path between i and j in G_0 .
- The adversary replied 1 to all edges on this path.
- Let (i', j') be the edge on this path that was queried last. Consider the graph when the adversary receives this query.



- Even if the adversary answers 0 on the edge (i', j') , when it sets all unqueried edges (including (i, j)) to 1, the graph must be connected.
- So the adversary should have answered 0 on (i', j') , a contradiction!